



Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird

Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023





Model free Flappy Bird


Michele Alessi, Samuele D'Avenia, Elena Rivaroli

MSc Data Science and Scientific Computing

Reinforcement Learning A.Y. 2022-2023




Flappy bird



- Tap -> The bird flies up
- Do not tap -> The bird moves down

Aim: keep the bird alive as long as possible

Flappy bird



- Tap -> The bird flies up
- Do not tap -> The bird moves down

Aim: keep the bird alive as long as possible

Player	Best score
Michele	5
Samuele	29
Elena	6

Text-flappy-bird environment ^[1]


Rules:

- If you flap you move up by 1
- If you don't, you move down

Reward: +1 for every step until
the game ends

Game ends:

- It touches the pipe
- It touches the floor




Model-free approach

States (dx, dy):


- 14 possible dx
- 22 possible dy

Actions

- 0 Remain Idle
- 1 Flap



General setup




General setup

- Techniques tried out
 - Monte Carlo Control*
 - Sarsa*
 - Expected Sarsa*
 - Q-learning**

*On-policy technique

** Off-policy technique




General setup


- Techniques tried out
 - Monte Carlo Control*
 - Sarsa*
 - Expected Sarsa*
 - Q-learning**
- Study on behaviour of hyper-parameters
 - λ on Sarsa and Q-learning
 - k_α on Sarsa, Q-learning and Expected Sarsa
 - k_ϵ on all the techniques

*On-policy technique

** Off-policy technique



Monte Carlo Control



This is a on-policy control method based on first-visit MC


Update the current estimate of action-value function:

$$Q(S_t, A_t) \leftarrow \text{average}(\text{Returns}(S_t))$$

$\text{Returns}(S_t, A_t)$ stores the first-visit return of (S_t, A_t) for all the episodes generated so far.


MC control cumulative rewards

$$\epsilon_0 = 0.2, k = 0.0$$



SARSA: On-Policy TD Control

It consists of two main ideas:



- Update the current estimate of action-value function

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- Construct an ϵ -greedy policy $\pi_t^\epsilon(s)$ given the current Q-value

Expected SARSA


Very simple modification to SARSA method.

Update the current estimate of action-value function as follows:

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \mathbb{E}_{\pi}[Q(S_{t+1}, A_{t+1}) \mid S_{t+1}] - Q(S_t, A_t) \right] \\ &\leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right] \end{aligned}$$

Q-Learning: Off-Policy TD Control

Update the current estimate of action-value function:





$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

SARSA control cumulative rewards

Expected SARSA cumulative rewards

$$\epsilon_0 = 0.2, k_\epsilon = 0, \alpha_0 = 0.15, k_\alpha = 0$$



	SARSA	Expected SARSA
Mean	55.374	49.982
Standard Deviation	52.530	42.942

TD(λ) Learning

The eligibility trace for each state s at time t is denoted as $e_t(s)$

On each step, the eligibility traces are updated as follows:

$$e_t(s, a) \leftarrow \begin{cases} \gamma\lambda e_{t-1}(s, a) + 1 & \text{if } s = s_t, a = a_t \\ \gamma\lambda e_{t-1}(s) & \text{otherwise} \end{cases}$$

- SARSA(λ) \rightarrow Fairly straightforward
- Q(λ) \rightarrow Some care handling eligibility traces

Hyperparameters

All the hyperparameters:

- t^* : until t^* we keep α and ε constant
- λ : determines eligibility decay

- α_0 : initial value of α
- ε_0 : initial value of ε
- k_α : controls how fast α goes to 0
- k_ε : controls how fast ε goes to 0




Change these

Convergence & Exploitation


Use constant α and ϵ up to some point t^* , and then decrease them as

- $$\alpha(t) = \frac{\alpha_0}{1 + k_\alpha(t - t^*)^{0.75}}$$


- $$\epsilon(t) = \frac{\epsilon_0}{1 + k_\epsilon(t - t^*)^{1.05}}$$



$$\varepsilon_0 = 0.2, k_\varepsilon = 0, \alpha_0 = 0.15, k_\alpha = 0, \lambda = 0.0$$

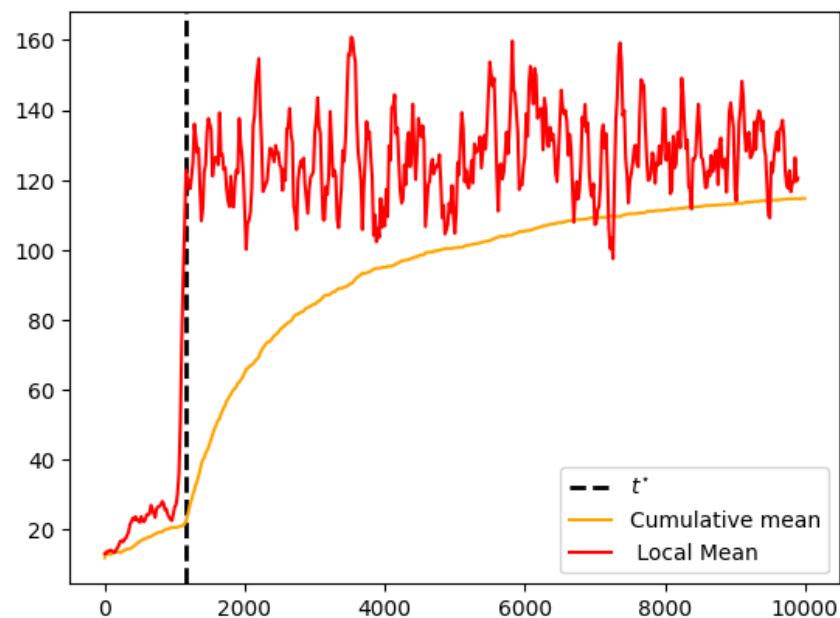


$$\varepsilon_0 = 0.2, k_\varepsilon = 5e-05, \alpha_0 = 0.15, k_\alpha = 3e-05, \lambda = 0.5$$



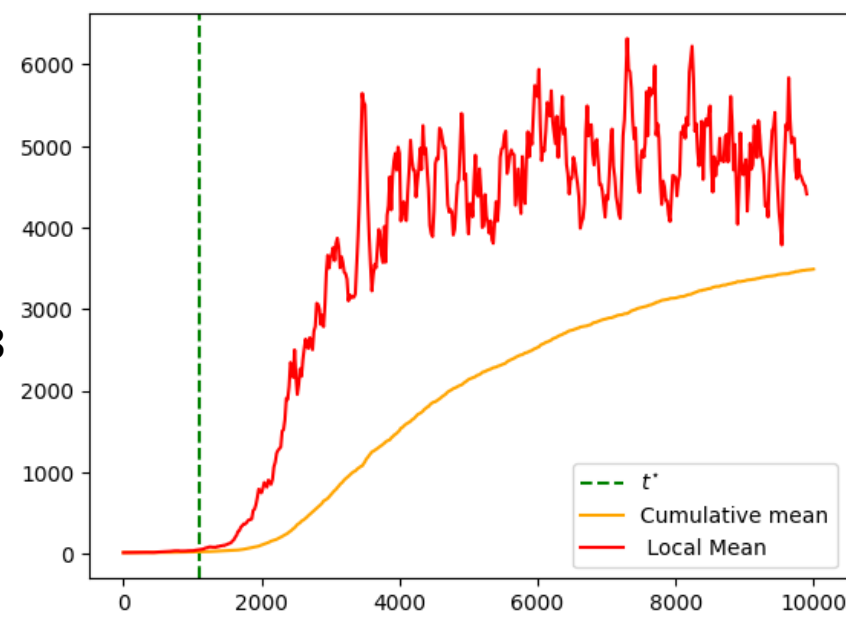
	Mean for fixed α and ε	Mean varying α and ε
Monte Carlo Control	48.74	199.46
SARSA	55.37	2633.70
Expected SARSA	49.98	127.30
Q-Learning	51.49	4783.61

Shown in
the plots




$k_\epsilon = 0.05$

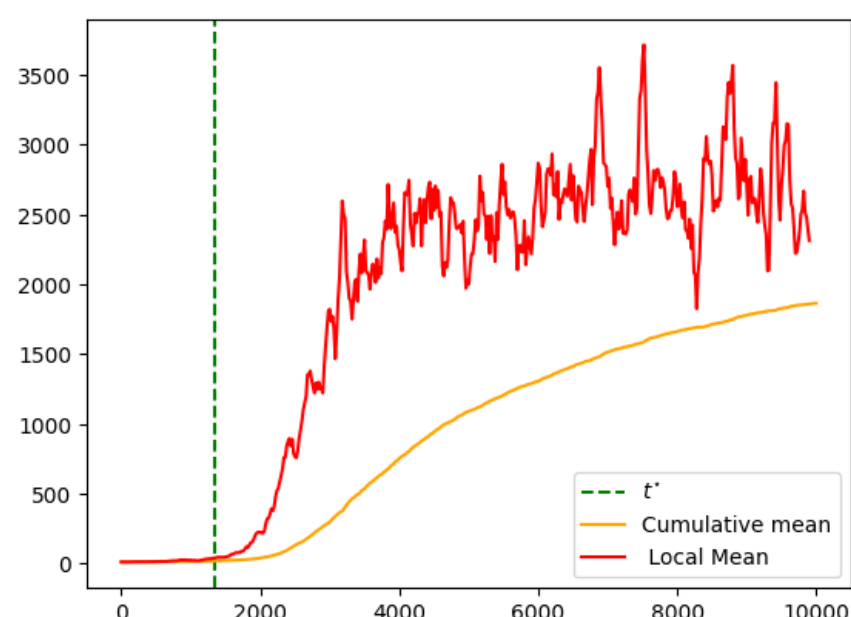
Q-Learning
 $\lambda = 0.5, k_\alpha = 0.00003$



$k_\epsilon = 0.00005$



SARSA
 $\lambda = 0.2, k_\alpha = 0.00003$




Final results

	Monte Carlo Control	SARSA	Expected SARSA	Q-Learning
λ	-	0.2	-	0.5
k_α	-	0.00003	0.0003	0.00003
k_ϵ	0.005	0.00005	0.00005	0.00005
Mean	199.46	2633.70	127.30	4783.61
Median	143	1831	93	3186
S.D.	175.70	2632.14	111.71	4952.15

Final results

	Monte Carlo Control	SARSA	Expected SARSA	Q-Learning
λ	-	0.2	-	0.5
k_α	-	0.00003	0.0003	0.00003
k_ϵ	0.005	0.00005	0.00005	0.00005
Mean	199.46	2633.70	127.30	4783.61
Median	143	1831	93	3186
S.D.	175.70	2632.14	111.71	4952.15




FLAP



IDLE



NOT VISITED



Further Work

- Perform a more in depth analysis of the hyper-parameters.
- More complex environment where agent receives frames as observations.

Thanks for the attention!



References

[1] <https://gitlab-research.centralesupelec.fr/stergios.christodoulidis/text-flappy-bird-gym/-/tree/master>

Some notes about the code:

- For the implementation of the algorithms we started from the code provided by the course tutor, Emanuele Panizon <https://www.ictp.it/member/emanuele-panizon>
- Small sections of the code were implemented using ChatGPT, more details on where it was used can be found on the GitHub repository.

Appendix: Pseudocodes

Pseudocodes were taken from

<http://incompleteideas.net/book/first/ebook/node1.html>

MC on policy control

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi \leftarrow$ an arbitrary ε -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$R \leftarrow$ return following the first occurrence of s, a

Append R to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$a^* \leftarrow \arg \max_a Q(s, a)$

For all $a \in \mathcal{A}(s)$:

$$\pi(s, a) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(s)| & \text{if } a = a^* \\ \varepsilon/|\mathcal{A}(s)| & \text{if } a \neq a^* \end{cases}$$

SARSA(λ)

Initialize $Q(s, a)$ arbitrarily and $e(s, a) = 0$, for all s, a

Repeat (for each episode):

Initialize s, a

Repeat (for each step of episode):

Take action a , observe r, s'

Choose a' from s' using policy derived from Q (e.g., ε -greedy)

$\delta \leftarrow r + \gamma Q(s', a') - Q(s, a)$

$e(s, a) \leftarrow e(s, a) + 1$

For all s, a :

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta e(s, a)$

$e(s, a) \leftarrow \gamma \lambda e(s, a)$

$s \leftarrow s'; a \leftarrow a'$

until s is terminal

$Q(\lambda)$

Initialize $Q(s, a)$ arbitrarily and $e(s, a) = 0$, for all s, a

Repeat (for each episode):

Initialize s, a

Repeat (for each step of episode):

Take action a , observe r, s'

Choose a' from s' using policy derived from Q (e.g., ε -greedy)

$a^* \leftarrow \arg \max_b Q(s', b)$ (if a' ties for the max, then $a^* \leftarrow a'$)

$\delta \leftarrow r + \gamma Q(s', a^*) - Q(s, a)$

$e(s, a) \leftarrow e(s, a) + 1$

For all s, a :

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta e(s, a)$

If $a' = a^*$, then $e(s, a) \leftarrow \gamma \lambda e(s, a)$

else $e(s, a) \leftarrow 0$

$s \leftarrow s'; a \leftarrow a'$

until s is terminal